



We use optional cookies to improve your experience on our websites, such as through social media connections, and to display personalized advertising based on your online activity. If you reject optional cookies, only cookies necessary to provide you the services will be used. You may change your selection by clicking "Manage Cookies" at the bottom of the page.
[Privacy Statement](#)
[Third-Party Cookies](#)

Accept

Reject

Manage cookies

GraphRAG Indexing

The GraphRAG indexing package is a data pipeline and transformation suite that is designed to extract meaningful, structured data from unstructured text using LLMs.

Indexing Pipelines are configurable. They are composed of workflows, standard and custom steps, prompt templates, and input/output adapters. Our standard pipeline is designed to:

- extract entities, relationships and claims from raw text

- perform community detection in entities
- generate community summaries and reports at multiple levels of granularity
- embed text into a vector space

The outputs of the pipeline are stored as Parquet tables by default, and embeddings are written to your configured vector store.

Getting Started

Requirements

See the [requirements](#) section in [Get Started](#) for details on setting up a development environment.

To configure GraphRAG, see the [configuration](#) documentation. After you have a config file you can run the pipeline using the CLI or the Python API.

Usage

CLI

```
uv run poe index --root <data_root> # default config mode
```



Python API

Please see the indexing API [python file](#) for the recommended method to call directly from Python code.

Further Reading

- To start developing within the *GraphRAG* project, see [getting started](#)
- To understand the underlying concepts and execution model of the indexing library, see [the architecture documentation](#)
- To read more about configuring the indexing engine, see [the configuration documentation](#)